

למידה סטטיסטית מבוססת נתונים - חורף - 096411

HW1

מגישים:

איתי ברקוביץ 039632732

אילן פרנק 043493386

הצורה הכללית של משוואת רגרסיה
 (2) $\sum (y_i - \hat{\beta}_0 - \hat{\beta}_1 x_i) \cdot x_i = 0$

$$\textcircled{1} \sum y_i - n \hat{\beta}_0 - \hat{\beta}_1 \sum x_i = 0$$

$$\boxed{\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}}$$

הצורה הכללית של משוואת רגרסיה

הצורה הכללית של משוואת רגרסיה

$$\textcircled{2} \sum (y_i - (\bar{y} - \hat{\beta}_1 \bar{x}) - \hat{\beta}_1 x_i) \cdot x_i = 0$$

$$= \sum \left[(y_i - \bar{y}) - \hat{\beta}_1 (x_i - \bar{x}) \right] x_i = \sum x_i (y_i - \bar{y}) - \hat{\beta}_1 \sum (x_i - \bar{x}) x_i$$

$$\Rightarrow \boxed{\hat{\beta}_1 = \frac{\sum_{i=1}^n (y_i - \bar{y}) x_i}{\sum_{i=1}^n (x_i - \bar{x}) x_i}}$$

הצורה הכללית של משוואת רגרסיה

$$\textcircled{1} \sum (y_i - \bar{y}) \cdot x_i = \sum (y_i - \bar{y})(x_i - \bar{x}) = \sum x_i y_i - n \bar{x} \bar{y}$$

$$\textcircled{2} \sum (x_i - \bar{x}) \cdot x_i = \sum (x_i - \bar{x})^2 = \sum x_i^2 - n \bar{x}^2$$

$$\hat{\beta}_1 = \frac{\sum x_i y_i - n \bar{x} \bar{y}}{\sum x_i^2 - n \bar{x}^2}$$

$$\bar{y} = \frac{1}{n} \sum y_i, \quad \bar{x} = \frac{1}{n} \sum x_i$$

הצורה הכללית של משוואת רגרסיה

$$\hat{\beta}_1 = \frac{\sum x_i y_i - \frac{1}{n} n \sum y_i \cdot \sum x_i}{\sum x_i^2 - \frac{1}{n} n (\sum x_i)^2}$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x} = \frac{\sum x_i y_i - \frac{1}{n} n \sum y_i \sum x_i}{\sum x_i^2 - \frac{1}{n} n (\sum x_i)^2} \times \frac{1}{n} \sum x_i$$

$$E\{\hat{\beta}_1\} = \beta_1$$

אם לא נזכר אחרת \hat{w}_1, \hat{w}_0 : צב.ד

$$E\{\hat{\beta}_0\} = \beta_0$$

$$\hat{\beta}_1 = \frac{\sum (x_i - \bar{x}) y_i}{S_{xx}} = \sum b_i y_i$$

$\underline{\underline{b_i}}$

המשקל

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x} = \frac{\sum y_i}{n} - \bar{x} \sum b_i y_i = \sum a_i y_i$$

$$\underline{\underline{a_i}} = \frac{1}{n} - b_i \bar{x} = \frac{1}{n} - \frac{(x_i - \bar{x}) \bar{x}}{S_{xx}}$$

$$E(\hat{\beta}_1) = E\left[\sum b_i y_i\right] = \sum b_i E(y_i) = \sum b_i (\beta_0 + \beta_1 x_i) = \beta_0 \sum b_i + \beta_1 \sum x_i b_i$$

$$\ast \sum b_i = \sum \frac{(x_i - \bar{x})}{S_{xx}} = \frac{1}{S_{xx}} \sum (x_i - \bar{x}) = 0$$

$$\ast \sum b_i x_i = \sum \frac{(x_i - \bar{x}) x_i}{S_{xx}} = \frac{1}{S_{xx}} \sum (x_i - \bar{x})^2 = \frac{S_{xx}}{S_{xx}} = 1.0$$

$$\Rightarrow E(\hat{\beta}_1) = \beta_1 \quad \text{c.v.}$$

$$E(\hat{\beta}_0) = E\left[\sum a_i y_i\right] = \sum a_i E(y_i) = \sum a_i (\beta_0 + \beta_1 x_i) = \beta_0 \sum a_i + \beta_1 \sum x_i a_i$$

$$\ast \sum a_i = \sum \left(\frac{1}{n} - \frac{(x_i - \bar{x}) \bar{x}}{S_{xx}} \right) = 1 - \bar{x} \sum \frac{(x_i - \bar{x})}{S_{xx}} = \underline{\underline{1.0}}$$

$$\ast \sum x_i a_i = \sum \left(\frac{1}{n} - \frac{(x_i - \bar{x}) \bar{x}}{S_{xx}} \right) \cdot x_i = \sum \frac{x_i}{n} - \bar{x} \sum \frac{(x_i - \bar{x}) \cdot x_i}{S_{xx}}$$

$$= \bar{x} - \bar{x} = 0$$

$$\Rightarrow E(\hat{\beta}_0) = \beta_0 \quad \text{c.v.}$$

In []:

In []:

```
# Question 2
#A.
import pandas as pd
import numpy as np

df = pd.read_csv('parkinsons_updrs_data.csv')

df['first'] = 1

df.head()
```

In [7]:

```
#B.
#Age -36 הנובדק הצעיר ביותר בן 85, הנובדק המבוגר ביותר בן 65, הנובדק גיל הנובדקים ממוצע גיל הנובדקים.
הערך מציג את גיל הנובדקים. ממוצע גיל הנובדקים הינו 65, הנובדק המבוגר ביותר בן 85, הנובדק הצעיר ביותר -36
בן
#sex מייצג את מין הנובדק. הערך הינו ערך בוליאני. 0 מייצג גבר ו-1 מייצג אישה. 32% מהנובדקים הינם גברים-1
הערך מייצג את מין הנובדק. הערך הינו ערך בוליאני. 0 מייצג גבר ו-1 מייצג אישה. 32% מהנובדקים הינם גברים-1
```

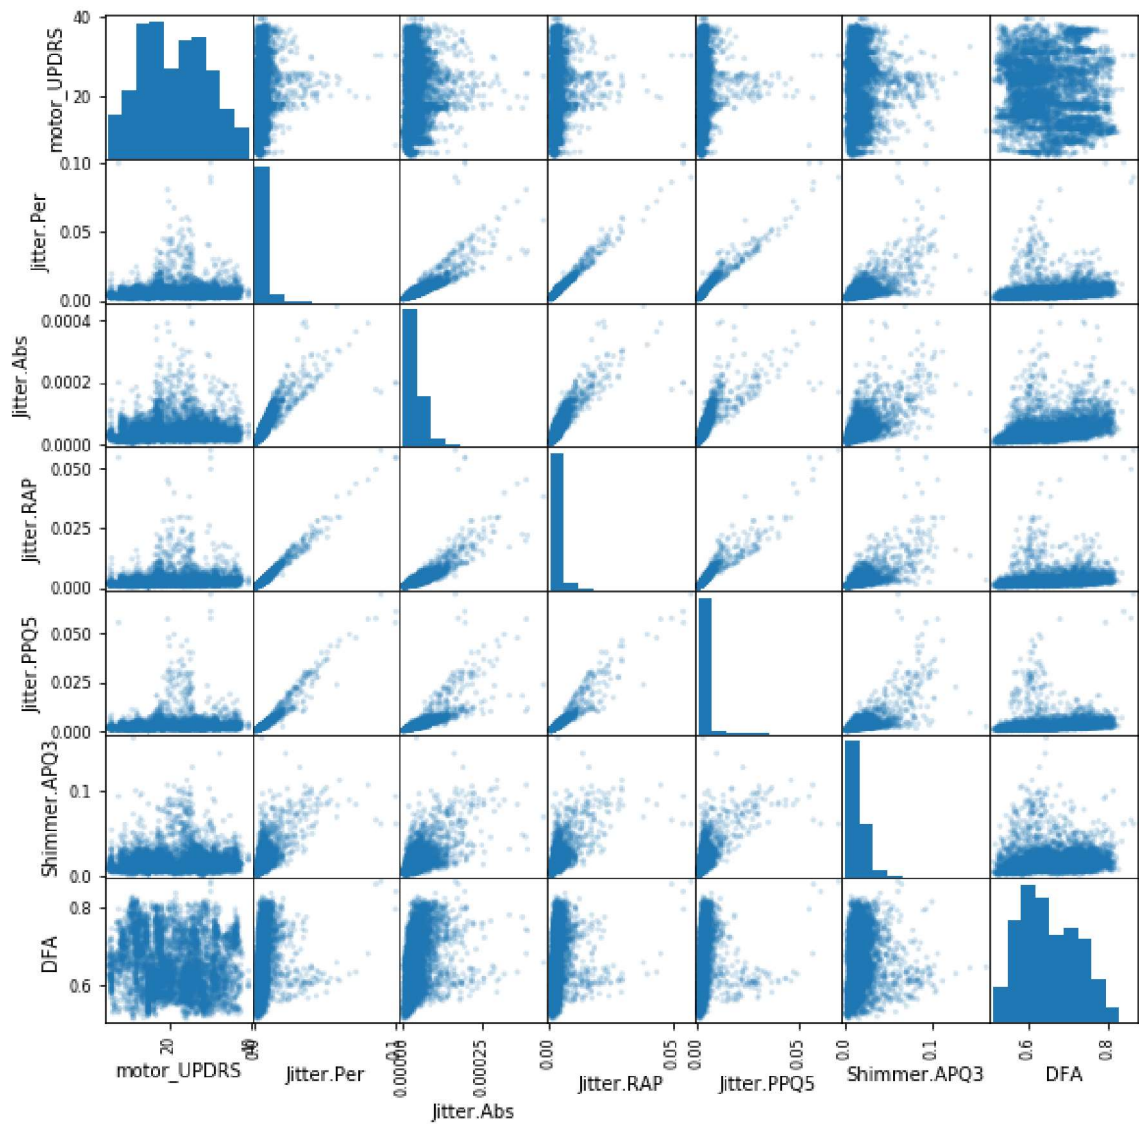
In [10]:

```
#C:
```

```
import numpy as np
import matplotlib.pyplot as plt
import pandas as pd
from sklearn import datasets
%matplotlib inline
```

```
df1 = df[['motor_UPDRS', 'Jitter.Per', 'Jitter.Abs', 'Jitter.RAP', 'Jitter.PPQ5', 'Shimmer.APQ3', 'DFA']]
pd.scatter_matrix(df1, alpha=0.2, figsize=(10, 10))
plt.show()
```

```
C:\Users\lenovo\Anaconda3\lib\site-packages\ipykernel_launcher.py:11: Futu
reWarning: pandas.scatter_matrix is deprecated, use pandas.plotting.scatte
r_matrix instead
# This is added back by InteractiveShellApp.init_path()
```



In [11]:

```
#D:

def answer_D(x,y):
    first_step = np.linalg.inv(np.dot(x.T,x))
    sec_step = np.dot(first_step,x.T)
    third_step = np.dot(sec_step,y)
    return third_step
```

In [12]:

```
#E:

x = np.array(df[['first','Jitter.Per', 'Jitter.Abs','Jitter.RAP','Jitter.PPQ5', 'Shimmer.APQ3', 'DFA']])
y = np.array(df[['motor_UPDRS']])
w = answer_D(x,y)
w
```

Out[12]:

```
array([[ 3.01179617e+01],
       [ 1.35025189e+03],
       [-2.05976226e+04],
       [-1.19055662e+03],
       [-7.51684784e+02],
       [ 4.34744789e+01],
       [-1.67630228e+01]])
```


In [18]:

```
#F:
import os
import pandas as pd
from statsmodels.formula.api import ols

df2 = df[['motor_UPDRS', 'first', 'Jitter.Per', 'Jitter.Abs', 'Jitter.RAP', 'Jitter.PPQ5',
          'Shimmer.APQ3', 'DFA']]
df2['motor_UPDRS'].values.reshape(-1, 1)
df2 = df2.rename(index=str, columns={'Jitter.Per': 'Jitter_Per', 'Jitter.Abs': 'Jitter_Abs',
                                     'Jitter.RAP': 'Jitter_RAP', 'Jitter.PPQ5': 'Jitter_PPQ5',
                                     'Shimmer.APQ3': 'Shimmer_APQ3'})
est = ols(formula = ' motor_UPDRS ~ first + Jitter_Per + Jitter_Abs + Jitter_RAP + Jitter_PPQ5 + Shimmer_APQ3 + Shimmer_APQ3 + DFA', data = df2).fit()
est.summary()
```

Out[18]:

OLS Regression Results

Dep. Variable:	motor_UPDRS	R-squared:	0.038			
Model:	OLS	Adj. R-squared:	0.037			
Method:	Least Squares	F-statistic:	38.61			
Date:	Wed, 14 Nov 2018	Prob (F-statistic):	2.94e-46			
Time:	08:47:22	Log-Likelihood:	-20533.			
No. Observations:	5875	AIC:	4.108e+04			
Df Residuals:	5868	BIC:	4.113e+04			
Df Model:	6					
Covariance Type:	nonrobust					
	coef	std err	t	P> t 	[0.025	0.975]
Intercept	15.0590	0.509	29.600	0.000	14.062	16.056
first	15.0590	0.509	29.600	0.000	14.062	16.056
Jitter_Per	1350.2519	158.081	8.542	0.000	1040.355	1660.148
Jitter_Abs	-2.06e+04	6726.924	-3.062	0.002	-3.38e+04	-7410.373
Jitter_RAP	-1190.5566	191.788	-6.208	0.000	-1566.532	-814.582
Jitter_PPQ5	-751.6848	126.651	-5.935	0.000	-999.967	-503.403
Shimmer_APQ3	43.4745	10.981	3.959	0.000	21.947	65.002
DFA	-16.7630	1.598	-10.489	0.000	-19.896	-13.630
Omnibus:	476.923	Durbin-Watson:	0.082			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	171.811			
Skew:	0.150	Prob(JB):	4.92e-38			
Kurtosis:	2.218	Cond. No.	2.36e+17			

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 2.55e-31. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

In []:

```
#f: בשני המקרים קיבלנו את אותם הערכים לאמד
```

סעיף ז:

עבור כללי ערכי האמדן קיבלנו ערכי p-value נמוכים מ-0.01, לכן נוכל לדחות את השארת האפס על β_j ברמת מובהקות של $\alpha = 0.01$.

ע"מ לדחות השארה זו נשתמש במבחן t דו צדדי בעל הסטטיסטי הבא:

$$\frac{\hat{\beta}_1}{\sqrt{\widehat{Var}(\hat{\beta}_1)}}$$

עבור רמת מובהקות $\alpha = 0.001$, נוכל לדחות את השארת האפס על β_j רק עבור המשתנה jitter abs, מכיוון שערך ה-p-value הינו 0.002.

$$p_r(y_i=1|x_i) = \frac{e^{\langle w, x_i \rangle}}{1 + e^{\langle w, x_i \rangle}} = \frac{1}{e^{-\langle w, x_i \rangle} + 1} \quad y \in \{+1\} \quad (3)$$

$$p_r(y_i=-1|x_i) = 1 - \frac{1}{e^{-\langle w, x_i \rangle} + 1} = \frac{e^{-\langle w, x_i \rangle}}{e^{-\langle w, x_i \rangle} + 1} = \frac{1}{1 + e^{\langle w, x_i \rangle}}$$

$$L(w) = \prod_{i=1}^m \left[\underbrace{p_r(y_i=1|x_i)}_{y_i=1 \text{ עבור } y_i=1}^{\frac{y_i+1}{2}} \cdot \underbrace{p_r(y_i=-1|x_i)}_{y_i=-1 \text{ עבור } y_i=-1}^{\frac{-(y_i-1)}{2}} \right] \quad \text{הפונקציה}$$

ועל ידי גזירה לאחור הפונקציה:

$$l(w) = \sum_{i=1}^m \left[\frac{y_i+1}{2} \log \frac{1}{1+e^{-\langle w, x_i \rangle}} + \frac{y_i-1}{2} \log \frac{1}{1+e^{\langle w, x_i \rangle}} \right]$$

נחפש w שממקסם דרך זו $\arg \max_w l(w)$

ה. נגזיר דבריו כנגד הסכום עבור צירי y השונים:

$$\begin{aligned} \text{עבור } y_i=1: & \quad \frac{2}{2} \log \frac{1}{1+e^{-\langle w, x_i \rangle}} + 0 \cdot \log \frac{1}{1+e^{\langle w, x_i \rangle}} \\ & = \log(1+e^{-\langle w, x_i \rangle})^{-1} \\ & = -\log(1+e^{-y_i \langle w, x_i \rangle}) \quad y_i=1 \end{aligned}$$

$$\begin{aligned} \text{עבור } y_i=-1: & \quad 0 \cdot \log \frac{1}{1+e^{-\langle w, x_i \rangle}} + \frac{-2}{2} \log \frac{1}{1+e^{\langle w, x_i \rangle}} \\ & = \log(1+e^{\langle w, x_i \rangle})^{-1} \quad y_i=-1 \\ & = -\log(1+e^{-y_i \langle w, x_i \rangle}) \end{aligned}$$

כאשר ניקח את כל $l(w)$ כסכום

$$l(w) = \sum_{i=1}^m -\log(1+e^{-y_i \langle w, x_i \rangle})$$

$$\begin{aligned} \arg \max_w l(w) &= \arg \max_w -\sum \log(1+e^{-y_i \langle w, x_i \rangle}) \quad \text{ועל ידי} \\ &= \arg \min_w \sum \log(1+e^{-y_i \langle w, x_i \rangle}) \end{aligned}$$

כנגד

In []:

Question 4

```

from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
import pandas as pd
import numpy as np

```

In [2]:

```

# (a)
iris = datasets.load_iris()
X = iris.data
y = iris.target

```

In [3]:

```

# (b)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=
1000)

```

In [4]:

```

# (c+d)
y_one_vs_all_train = {}
y_one_vs_all_test = {}
lrs = {}
print(y_test)
for i in range(3):
    y_one_vs_all_train[i] = np.ones(y_train.shape)*-1
    y_one_vs_all_train[i][y_train==i] = 1

    y_one_vs_all_test[i] = np.ones(y_test.shape)*-1
    y_one_vs_all_test[i][y_test==i] = 1
#     print(y_one_vs_all_test[i])

    lrs[i] = LogisticRegression()
    lrs[i].fit(X_train,y_one_vs_all_train[i])
#     print(lrs[i].predict_proba(X_test))

    print(f'{i} score: {lrs[i].score(X_test, y_one_vs_all_test[i])}')
# y_one_vs_all_train, y_train

```

```

[1 0 2 2 0 0 1 1 0 2 2 1 0 0 2 1 2 1 0 0 1 1 2 2 1 0 1 1 1 2 0 1 2 1 0 2 2
 1]
0 score: 1.0
1 score: 0.6842105263157895
2 score: 0.9473684210526315

```


In [7]:

```
# (e)

def get_one_vs_all(lrs, test):
    predictions = [max(lrs.keys(), key=lambda i: lrs[i].predict_proba([x])[0, 1]) for x
in test]
    # for sample_idx, sample in enumerate(test):
    #     print([lrs[i].predict_proba([sample]) for i in lrs])
    #     max_lr = max(lrs.keys(), key=lambda i: lrs[i].predict_proba([sample])[0, 1])
    #     print(max_lr)
    return predictions

y = get_one_vs_all(lrs, X_test)

y == y_test
```

Out[7]:

```
array([ True,  True,  True,  True,  True,  True, False,  True,  True,
        True,  True,  True,  True,  True,  True,  True,  True,  True,
        True,  True,  True, False,  True,  True,  True,  True, False,
        True,  True,  True,  True,  True,  True,  True,  True,  True,
        True, False])
```

In [397]:

```

# (f)

from sklearn.metrics import confusion_matrix
import numpy as np
import matplotlib.pyplot as plt
from sklearn.metrics import confusion_matrix
import itertools

def plot_confusion_matrix(cm, classes,
                          normalize=False,
                          title='Confusion matrix',
                          cmap=plt.cm.Blues):
    """
    This function prints and plots the confusion matrix.
    Normalization can be applied by setting `normalize=True`.
    """
    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation=45)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.ylabel('True label')
    plt.xlabel('Predicted label')
    plt.tight_layout()

y_pred = get_one_vs_all(lrs, X_test)
# print(y_test)
# print(y_pred)
print(list(zip(y_test, y_pred)))
cnf_matrix = confusion_matrix(y_test, y_pred)
np.set_printoptions(precision=2)



# Plot non-normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=[0,1],
                      title='Confusion matrix, without normalization')

# Plot normalized confusion matrix
plt.figure()
plot_confusion_matrix(cnf_matrix, classes=[0,1], normalize=True,
                      title='Normalized confusion matrix')

plt.show()

```

```
[(1, 1), (0, 0), (2, 2), (2, 2), (0, 0), (0, 0), (1, 2), (1, 1), (0, 0),  
(2, 2), (2, 2), (1, 1), (0, 0), (0, 0), (2, 2), (1, 1), (2, 2), (1, 1),  
(0, 0), (0, 0), (1, 1), (1, 2), (2, 2), (2, 2), (1, 1), (0, 0), (1, 2),  
(1, 1), (1, 1), (2, 2), (0, 0), (1, 1), (2, 2), (1, 1), (0, 0), (2, 2),  
(2, 2), (1, 2)]
```



In [398]:

```

# (g)
# print(list(zip(X_train, y_train)))
for i, sample in enumerate(X_test):
    print(f'X: {sample}, true label: {y_test[i]}, classifier 1: {lrs[1].predict_proba([sample])[0, 1]}, '
          f'classifier 2: {lrs[2].predict_proba([sample])[0, 1]}')
print()
for i in range(3):
    print(f'classifier {i}:')
    print(lrs[i].coef_)
print()
print(f'average train sample for classifier 1: {np.average(X_train[y_train==1], axis=0)}')
print(f'average train sample or classifier 2: {np.average(X_train[y_train==2], axis=0)}')
print()
print(f'average test for classifier 1: {np.average(X_test[y_test==1], axis=0)}')
print(f'average test for classifier 2: {np.average(X_test[y_test==2], axis=0)}')
print()
print("""The last test sample is labeled as 1, but tagged as 2 by the one-vs-all classifier.
We can tell by looking at the classifiers' weights that classifier 1 relies mostly on features 0,2, while 1,3 negate it,
and classifier 2 relies mostly on features 2,3, while 0,1 negate it.

We look at the data for the problematic sample: [ 5.6  3.  4.5  1.5].
Comparing to other samples of label 1, this has a relatively high value for feature 1 and 2.
This causes the score for classifier 1 be relatively low (0.27) and for classifier 2 be a little high (0.44),
which causes the sample to be labeled as 2.""")

```

X: [5.7 3. 4.2 1.2], true label: 1, classifier 1: 0.3184256754872217, classifier 2: 0.15081079174784146
X: [4.9 3.1 1.5 0.1], true label: 0, classifier 1: 0.23192467152550447, classifier 2: 0.00011700991022690092
X: [6.1 3. 4.9 1.8], true label: 2, classifier 1: 0.27960008167133626, classifier 2: 0.6287402302847547
X: [6.9 3.2 5.7 2.3], true label: 2, classifier 1: 0.2543561079263063, classifier 2: 0.8682121895240641
X: [4.8 3.4 1.9 0.2], true label: 0, classifier 1: 0.1485946800469467, classifier 2: 0.00030902795364945776
X: [5.2 3.4 1.4 0.2], true label: 0, classifier 1: 0.15519102366489604, classifier 2: 5.076884717174568e-05
X: [5.4 3. 4.5 1.5], true label: 1, classifier 1: 0.2342521975914924, classifier 2: 0.5336967442627625
X: [6.2 2.2 4.5 1.5], true label: 1, classifier 1: 0.671426481741066, classifier 2: 0.41278844180816826
X: [4.4 3. 1.3 0.2], true label: 0, classifier 1: 0.1778857122143906, classifier 2: 0.00024327253363154723
X: [6.3 2.7 4.9 1.8], true label: 2, classifier 1: 0.42641876195744016, classifier 2: 0.6256782301284715
X: [6.4 2.8 5.6 2.2], true label: 2, classifier 1: 0.352800488421773, classifier 2: 0.9380449220273951
X: [6.4 2.9 4.3 1.3], true label: 1, classifier 1: 0.4428950199101907, classifier 2: 0.08659259870469917
X: [5.1 3.8 1.9 0.4], true label: 0, classifier 1: 0.07608344215127996, classifier 2: 0.00018888398064920722
X: [5. 3.6 1.4 0.2], true label: 0, classifier 1: 0.10254802431707455, classifier 2: 5.7347928824759955e-05
X: [6.7 3.3 5.7 2.5], true label: 2, classifier 1: 0.16630785575791834, classifier 2: 0.9287465297500211
X: [5.8 2.7 4.1 1.], true label: 1, classifier 1: 0.5060438308245141, classifier 2: 0.09544381343775583
X: [6. 2.2 5. 1.5], true label: 2, classifier 1: 0.6873627491159601, classifier 2: 0.752979646735354
X: [5.8 2.7 3.9 1.2], true label: 1, classifier 1: 0.42756861283434566, classifier 2: 0.09550194938571142
X: [4.9 3.1 1.5 0.1], true label: 0, classifier 1: 0.23192467152550447, classifier 2: 0.00011700991022690092
X: [5. 3.3 1.4 0.2], true label: 0, classifier 1: 0.161969454587434, classifier 2: 7.947230466700494e-05
X: [5.5 2.4 3.7 1.], true label: 1, classifier 1: 0.5511030657944215, classifier 2: 0.08983671376181318
X: [5.9 3.2 4.8 1.8], true label: 1, classifier 1: 0.18835522010543446, classifier 2: 0.6042370158206046
X: [7.2 3.2 6. 1.8], true label: 2, classifier 1: 0.4555996065243735, classifier 2: 0.7157393298750881
X: [6.9 3.1 5.1 2.3], true label: 2, classifier 1: 0.24279935723380489, classifier 2: 0.6549465817014978
X: [4.9 2.4 3.3 1.], true label: 1, classifier 1: 0.41919438860707886, classifier 2: 0.09978971101872118
X: [5.7 4.4 1.5 0.4], true label: 0, classifier 1: 0.034461510064381765, classifier 2: 1.4417578310348812e-05
X: [6. 3.4 4.5 1.6], true label: 1, classifier 1: 0.16375716869686371, classifier 2: 0.2511714922251177
X: [6.2 2.9 4.3 1.3], true label: 1, classifier 1: 0.41245934172889936, classifier 2: 0.11747330256713163
X: [5. 2. 3.5 1.], true label: 1, classifier 1: 0.6261913982852128, classifier 2: 0.18495233800212407
X: [6.4 3.1 5.5 1.8], true label: 2, classifier 1: 0.33225893978418053, classifier 2: 0.7793908693880235
X: [5.1 3.3 1.7 0.5], true label: 0, classifier 1: 0.1396366180413344


```

5, classifier 2: 0.00025974398363687003
X: [ 6.   2.2  4.   1. ], true label: 1, classifier 1: 0.728103239958982,
classifier 2: 0.09363832901161794
X: [ 6.8  3.2  5.9  2.3], true label: 2, classifier 1: 0.2575614717406986,
classifier 2: 0.9245606022456869
X: [ 5.6  2.9  3.6  1.3], true label: 1, classifier 1: 0.2682431071545994,
classifier 2: 0.07062766125176291
X: [ 5.1  3.4  1.5  0.2], true label: 0, classifier 1: 0.1522405750784626
4, classifier 2: 7.537523022263318e-05
X: [ 6.7  2.5  5.8  1.8], true label: 2, classifier 1: 0.6588771943476974,
classifier 2: 0.8891610578909227
X: [ 6.4  2.7  5.3  1.9], true label: 2, classifier 1: 0.4514781835353867
7, classifier 2: 0.8133596518339681
X: [ 5.6  3.   4.5  1.5], true label: 1, classifier 1: 0.2572965193943098
6, classifier 2: 0.44907763659831124

```

```

classifier 0:
[[ 0.4   1.34 -2.1  -0.95]]
classifier 1:
[[ 0.62 -1.75  0.4  -1.18]]
classifier 2:
[[-1.7  -1.09  2.26  2.26]]

```

```

average train sample for classifier 1: [ 6.02  2.79  4.32  1.35]
average train sample or classifier 2: [ 6.59  2.99  5.58  2.03]

```

```

average test for classifier 1: [ 5.73  2.73  4.11  1.28]
average test for classifier 2: [ 6.57  2.92  5.45  2.   ]

```

The last test sample is labeled as 1, but tagged as 2 by the one-vs-all classifier.

We can tell by looking at the classifiers' weights that classifier 1 relies mostly on features 0,2, while 1,3 negate it, and classifier 2 relies mostly on features 2,3, while 0,1 negate it.

We look at the data for the problematic sample: [5.6 3. 4.5 1.5]. Comparing to other samples of label 1, this has a relatively high value for feature 1 and 2. This causes the score for classifier 1 be relatively low (0.27) and for classifier 2 be a little high (0.44), which causes the sample to be labeled as 2.